

THIS IS CS4048!

**GCR:wzj3vua**

# CONVOLUTION NEURAL NETWORKS

People telling me AI is going  
to destroy the world

My neural network




# NEURAL NETWORK WITH IMAGES

We know it is good to learn a small model.


From this fully connected model, do we really need all the edges?

Can some of these be shared?



An image is nothing but a matrix of pixel values, right? So why not just flatten the image (e.g. 3x3 image matrix into a 9x1 vector) and feed it to a Multi-Level Perceptron for classification purposes? Uh.. not really.

1	1	0
4	2	1
0	2	1




1
1
0
4
2
1
0
2
1

To capture Spatial/temporal dependencies.

# CNN

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.



1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

# Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Dot  
product

3

-1



# Convolution

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3

-3

# Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

# Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

Repeat this for each filter



Two 4 x 4 images  
Forming 2 x 4 x 4 matrix

# FOR GREY IMAGES



1	0.6	0.3	1	1
1	0	0	0.3	1
1	1	0	0.6	1
1	0.6	0	0	0.3
1	1	1	1	1

# FOR GREY IMAGES

1	0.6	0.3	1	1
1	0	0	0.3	1
1	1	0	0.6	1
1	0.6	0	0	0.3
1	1	1	1	1

0	1	0
1	2	1
0	1	0

# FOR GREY IMAGES

multiply numbers in same location

1 x 0	0.6 x 1	0.3 x 0	1	1
1 x 1	0 x 2	0 x 1	0.3	1
1 x 0	1 x 1	0 x 0	0.6	1
1	0.6	0	0	0.3
1	1	1	1	1

0 0.6 0


1 0 0

0 1 0

2.6

sum all

# FOR GREY IMAGES



# FOR GREY IMAGES

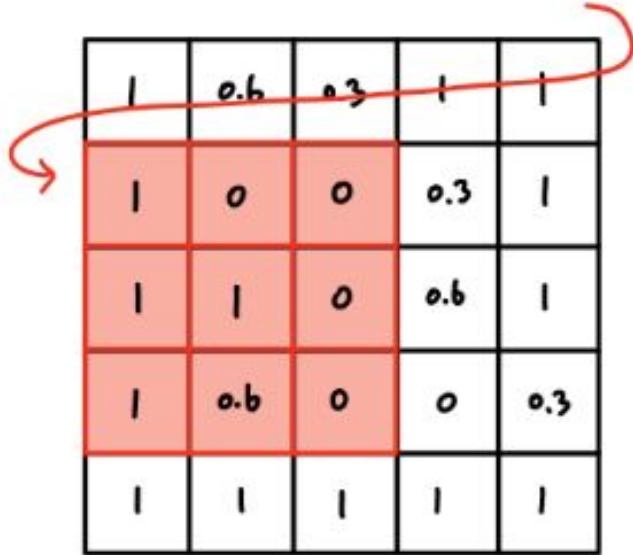
→

1	0.6	0.3	1	1
1	0	0	0.3	1
1	1	0	0.6	1
1	0.6	0	0	0.3
1	1	1	1	1

2.6	0.6	3.2
-----	-----	-----

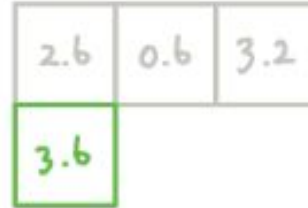


# FOR GREY IMAGES



A 5x5 grid of values. The first three rows and first three columns are highlighted in red. A red arrow points from the top-right corner of the red region (row 3, column 3) to the top-left corner of the red region (row 1, column 1), indicating a convolution operation.

1	0.6	0.3	1	1
1	0	0	0.3	1
1	1	0	0.6	1
1	0.6	0	0	0.3
1	1	1	1	1



A 2x3 grid of values. The bottom-left cell (row 2, column 1) is highlighted in green.

2.6	0.6	3.2
3.6		

# FOR GREY IMAGES

This sliding operation is called **convolution**.\*

filter

0	1	0
1	2	1
0	1	0

1	0.6	0.3	1	1
1	0	0	0.3	1
1	1	0	0.6	1
1	0.6	0	0	0.3
1	1	1	1	1


data

2.6	0.6	3.2
3.6	1.6	2.5
4.2	1.6	1.9

result


# FOR GREY IMAGES

filters




1	0.6	0.3	1	1
1	0	0	0.3	1
1	1	0	0.6	1
1	0.6	0	0	0.3
1	1	1	1	1

data



results


# FOR GREY IMAGES




# CNN

Result array from a convolution operation using a filter is called a **feature map**.

The number of feature maps is called the number of **channels** as well. If you use 10 filters, then there will be 10 output channels.




# Color image: RGB 3 channels



# CNN WITH RGB IMAGES

The filter moves to the right with a certain Stride Value till it parses the complete width. Moving on, it hops down to the beginning (left) of the image with the same Stride Value and repeats the process until the entire image is traversed.



0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...	...	...	...	...	...	...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...	...	...	...	...	...	...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...	...	...	...	...	...	...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

+

+

+ 1 = -25


↑  
Bias = 1

Output

-25				...
				...
				...
				...
...	...	...	...	...




# Convolution v.s. Fully Connected



Fully-  
connected

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0



# SO WHAT KIND OF VALUES SHOULD BE IN THE FILTERS?

Before the era of deep learning, the values in the filter were fixed. Have you ever used "sharpen filter" in a photo editor?

The values in a convolution filter are the weights of a Convolutional Neural Network.

filter (old)


0	1	0
1	2	1
0	1	0

fixed values


filter (deep learning)


$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

being optimized values  
(weights)




fewer parameters!





Fewer parameters

Even fewer parameters



# CNN

That means, we put some random values into the filters and set how many filters we use in each layer before starting training, but the values in the filters change as the training goes on.

By doing so, the values in the filters are getting optimized, more meaningful feature maps are being extracted, so the model makes better prediction.


# CNN

The objective of the Convolution Operation is to extract the high-level features such as edges, from the input image.


Conventionally, the first ConvLayer is responsible for capturing the Low-Level features such as edges, color, gradient orientation, etc.

With added layers, the architecture adapts to the High-Level features as well, giving us a network that has a wholesome understanding of images in the dataset, similar to how we would.

- Convolution layer takes an input feature map of dimension  $W \times H \times N$  and produces an output feature map of dimension  $\hat{W} \times \hat{H} \times M$
- Each layer is defined using following parameters:
  - # Input channels (N)
  - # Output channels (M)
  - Kernel size
  - Padding
  - Stride

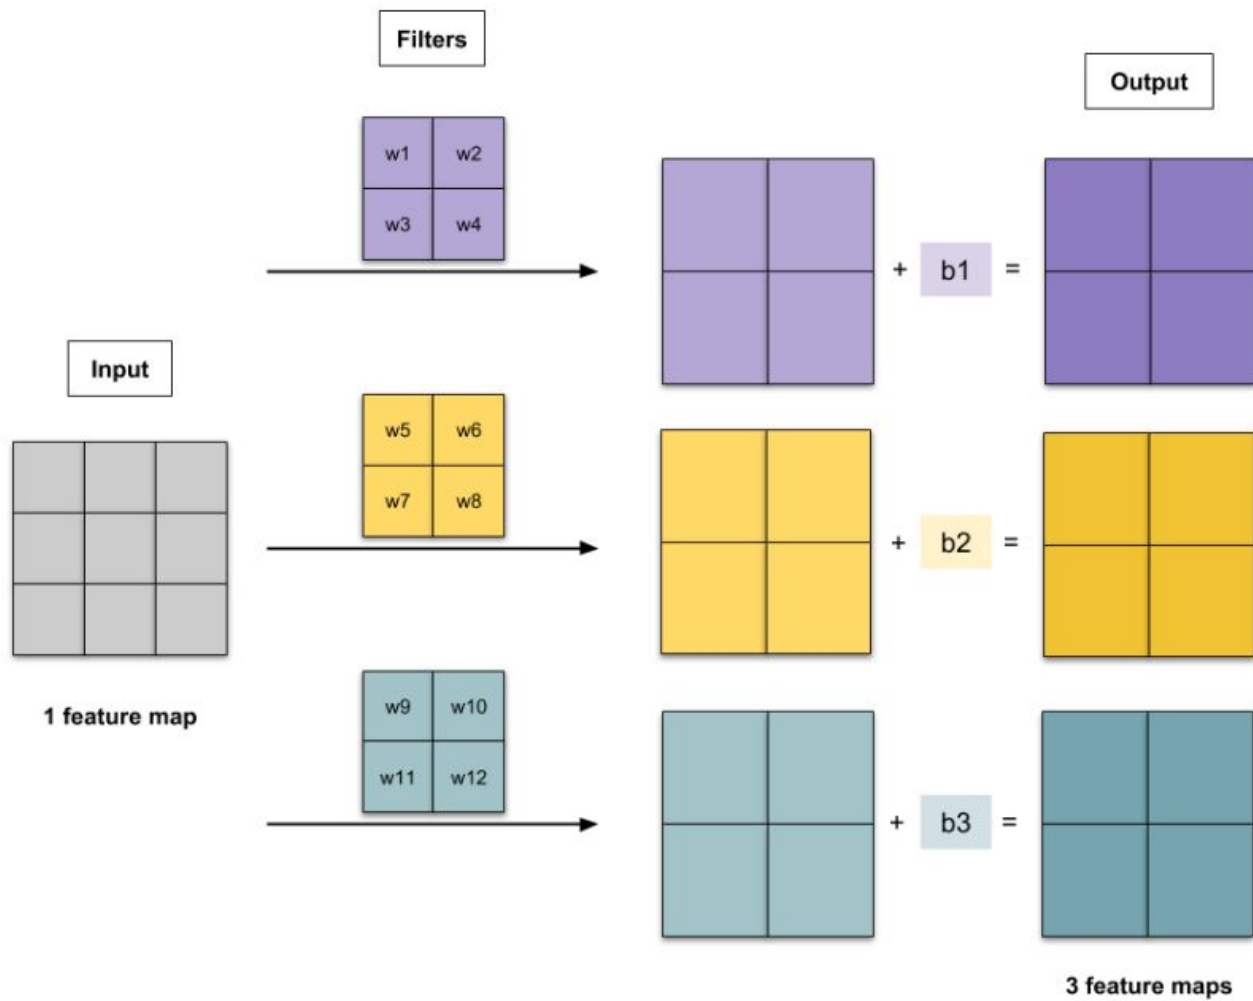


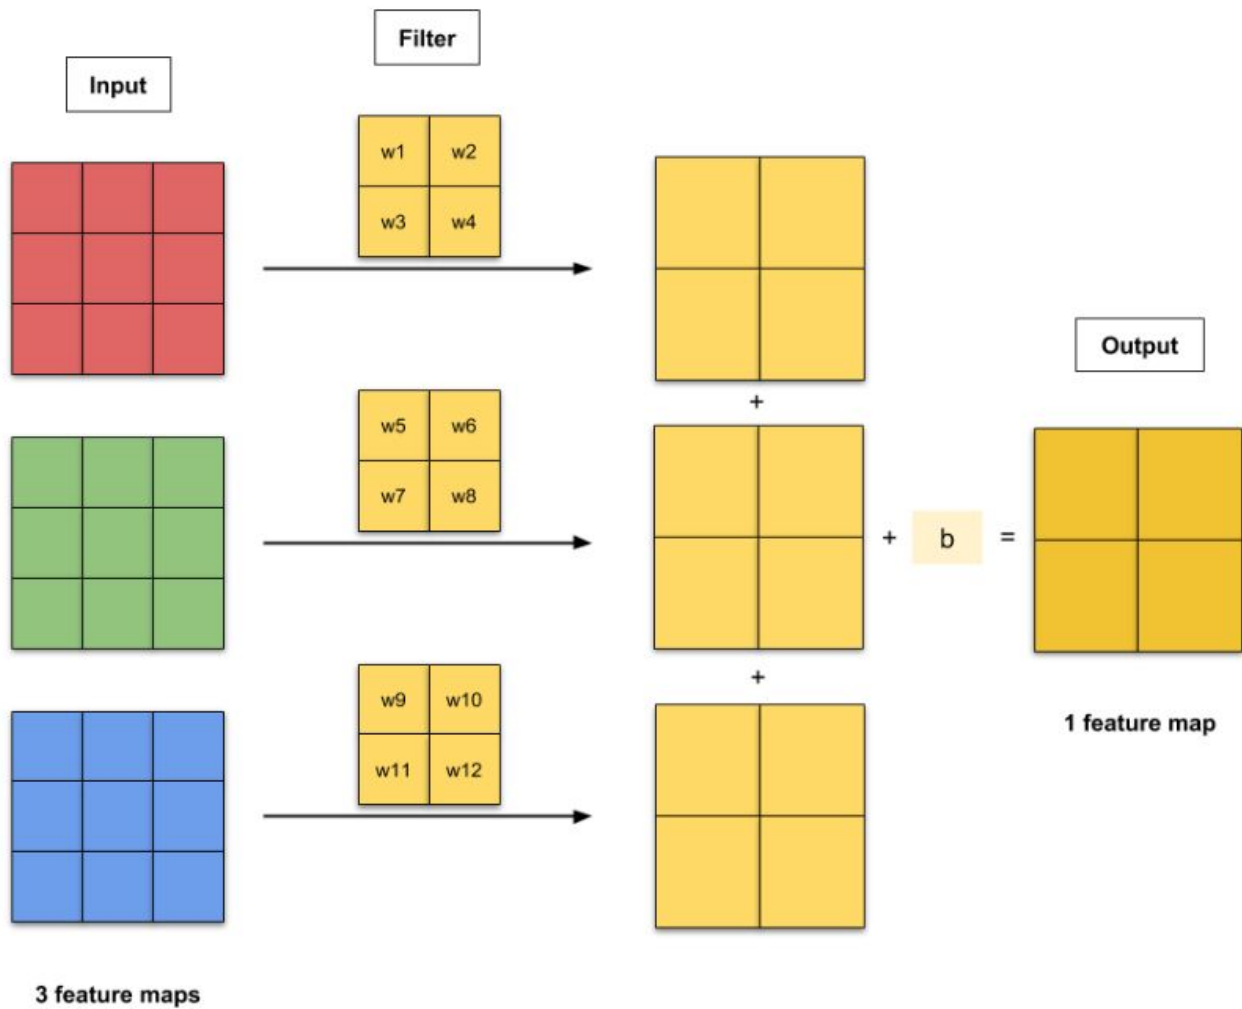
**Figure:** In this example, 5x5 input is convolved with 3x3 kernel with stride=padding=1 to produce an output of size 5x5.

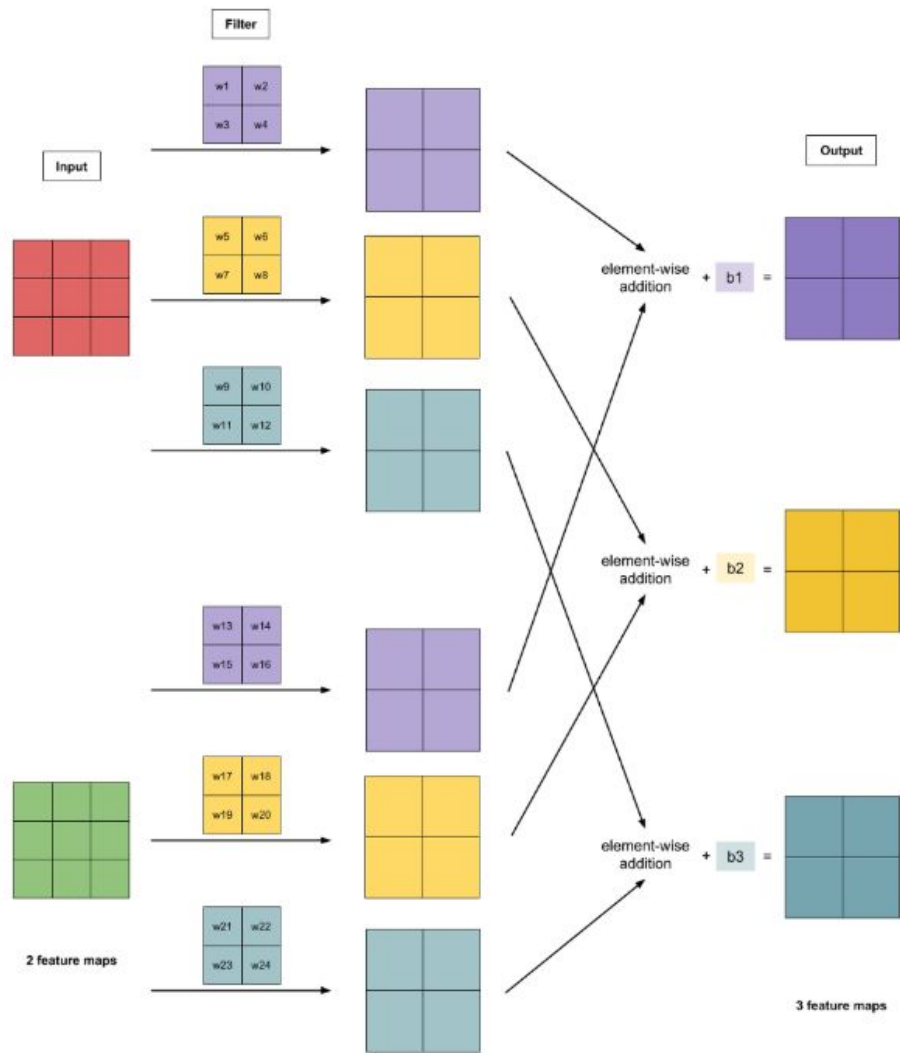


**Figure:** In this example, 5x5 input is convolved with 3x3 kernel with **stride=2** and padding=1 to produce an output of size 3x3.



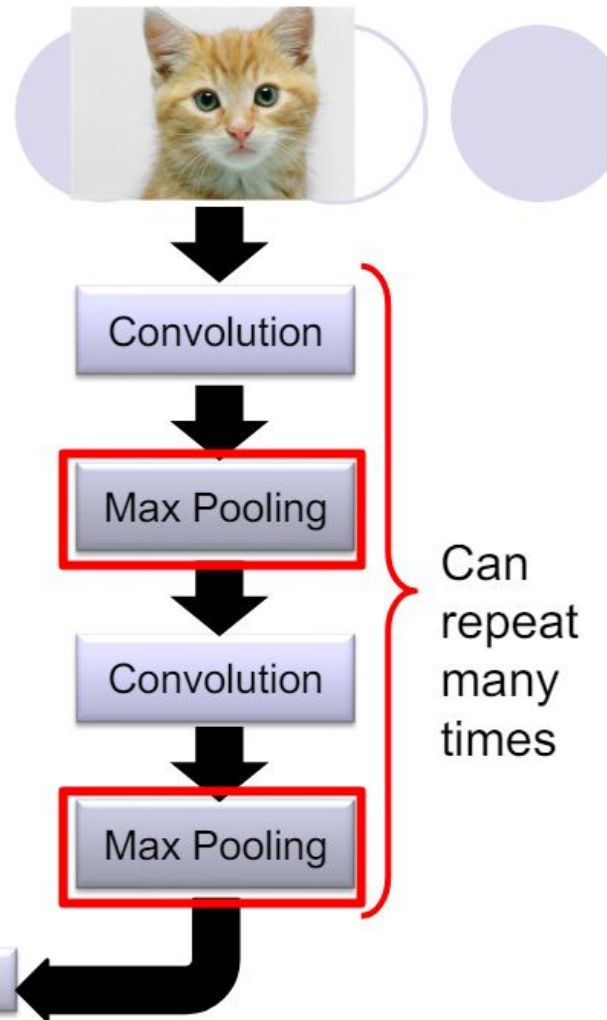







# The whole CNN

cat dog .....



# POOLING LAYERS


Similar to the Convolutional Layer, the Pooling layer is responsible for reducing the spatial size of the Convolved Feature.

Decreases the computational power required to process the data through dimensionality reduction

Extract dominant features.

Max - Min - Avg pooling

# MAX POOL



# MAX POOLING

Max Pooling also performs as a Noise Suppressant.

It discards the noisy activations altogether and also performs de-noising along with dimensionality reduction.

3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

# MAX POOLING

Pool size: 3x3, Stride: 1

-4	0	-2	4	1
3	1	0	2	1
1	0	1	1	1
4	6	5	1	0
-1	2	0	0	0

Features

Max Pooling

3	4
6	5

Min Pooling

-4	-2
-1	0

Average Pooling

0	1
2	1

Output



# Why Pooling

- Subsampling pixels will not change the object

bird



Subsampling

bird



We can subsample the pixels to make image

smaller



fewer parameters to characterize the image

# CNN

A CNN compresses a fully connected network in following ways:

- Reducing number of connections
- Shared weights on the edges
- Max pooling further reduces the complexity

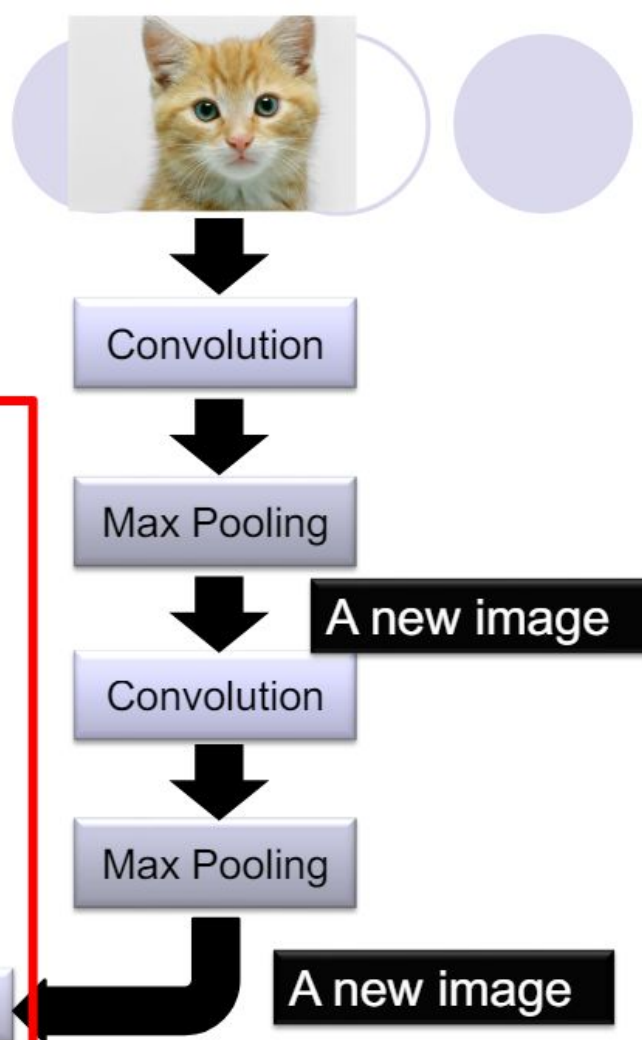

# CNN

Convolutional Neural Networks is extension of traditional Multi-layer Perceptron, based on 3 ideas:


- Local receive fields
- Shared weights
- Spatial / temporal sub-sampling

# The whole CNN


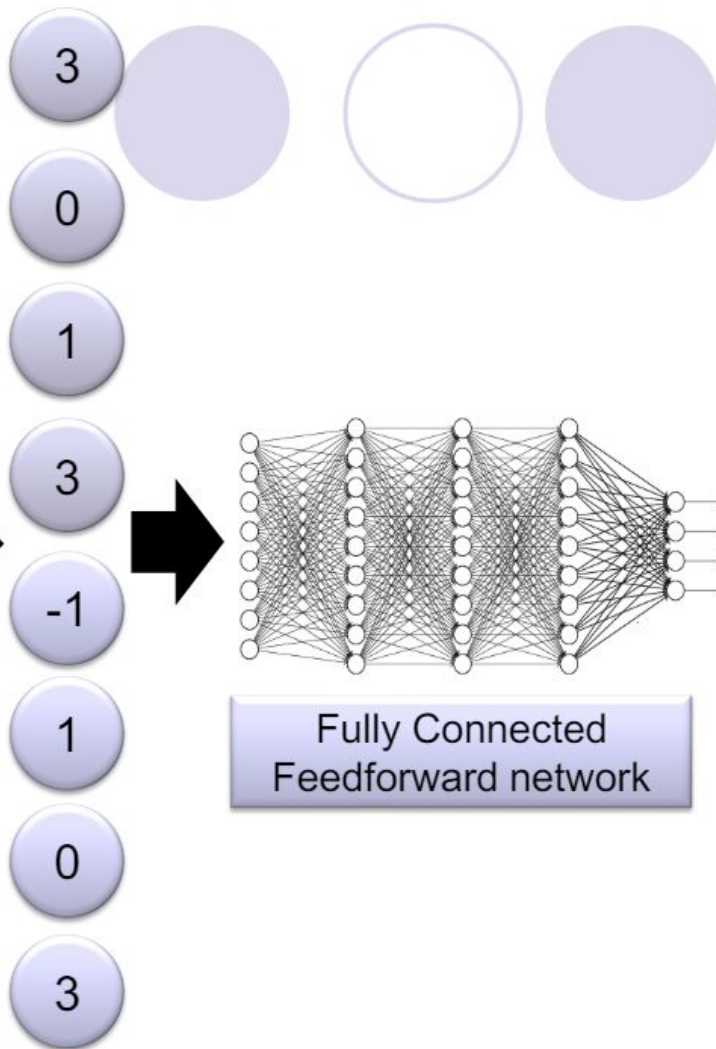
cat dog .....



# Flattening

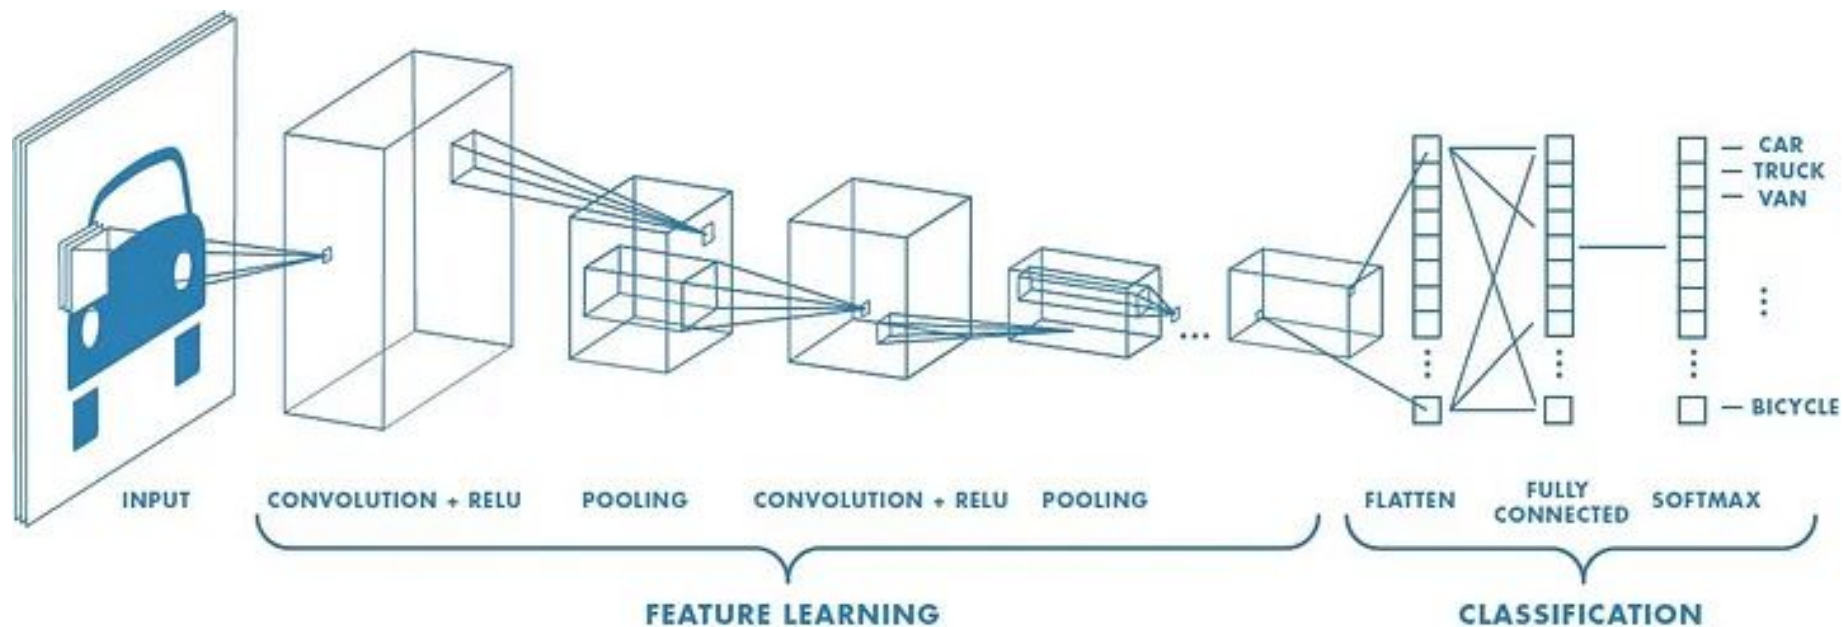


Flattened

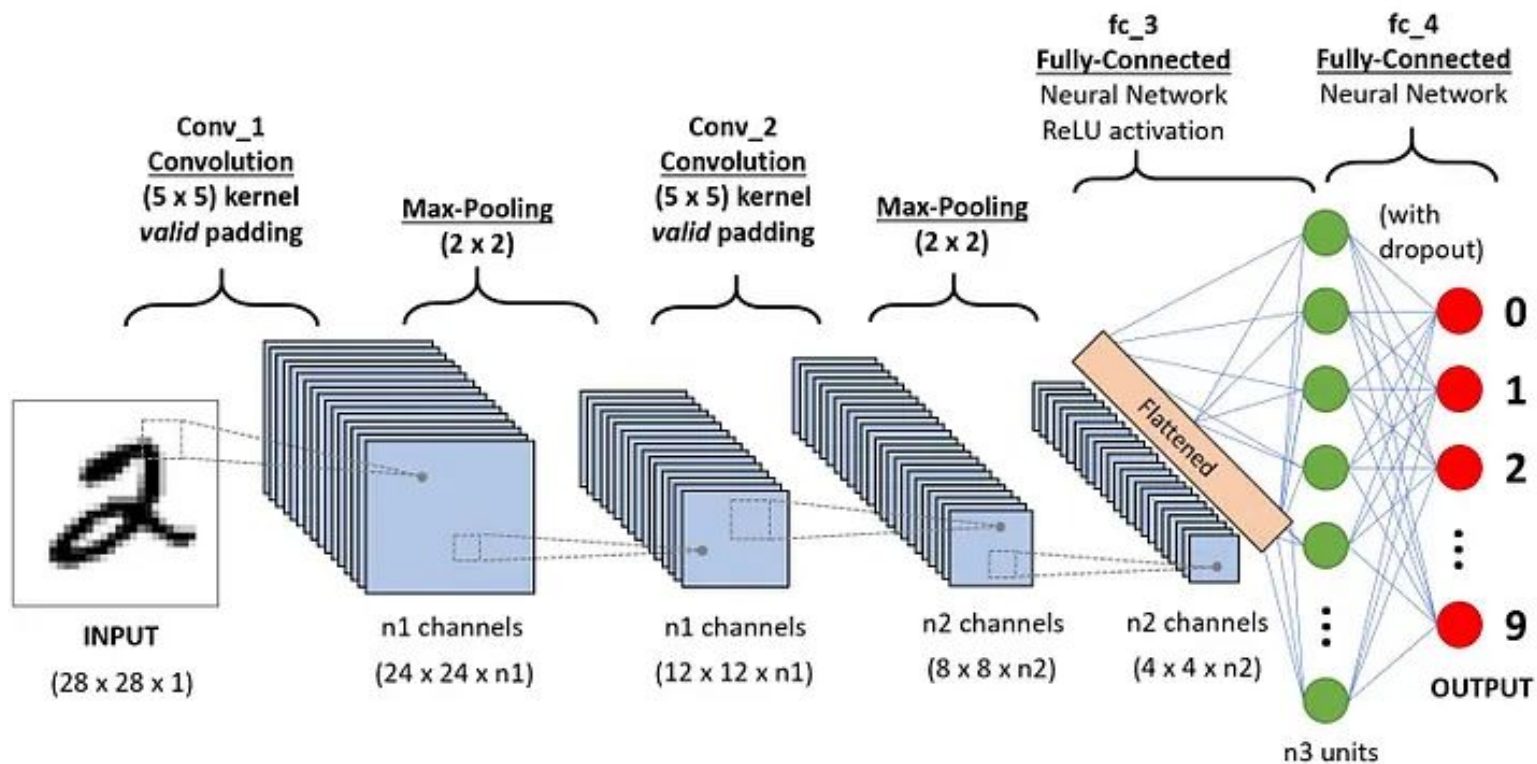


Fully Connected  
Feedforward network

# CNN



# CNN







**CONVOLUTIONAL NEURAL  
NETWORKS**



**IS THERE ANYTHING THEY CAN'T  
DO?**

[memegenerator.net](http://memegenerator.net)

# REFERENCES

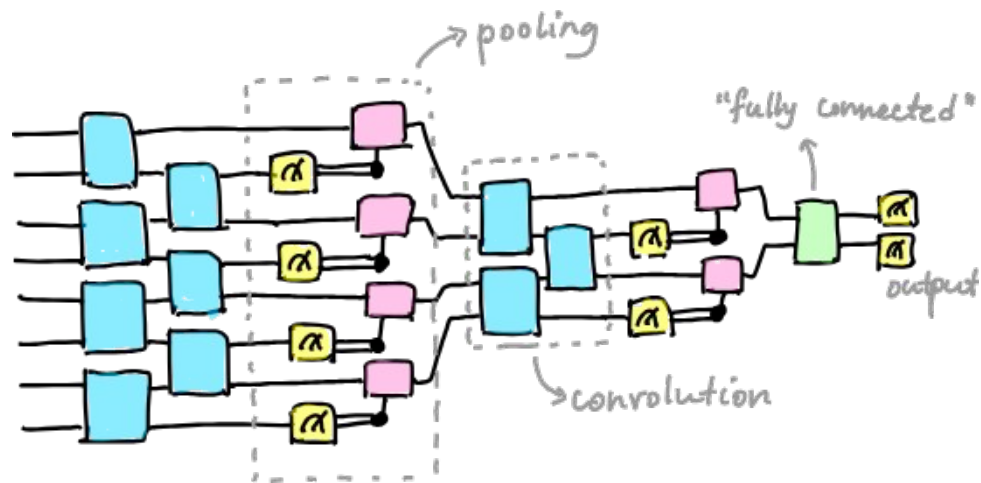
Some slides are taken from Dr Atif Tahir's slides.

<https://www.philgineer.com/2021/11/convolutional-neural-network-explained.html>

[https://github.com/vdumoulin/conv\\_arithmetic?source=post\\_page-----3bd2b1164a53-----](https://github.com/vdumoulin/conv_arithmetic?source=post_page-----3bd2b1164a53-----)

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

# QUANTUM CONVOLUTION NEURAL NETWORKS



quantum convolutional neural network



Nature

<https://www.nature.com> > nature physics > articles

## Quantum convolutional neural networks


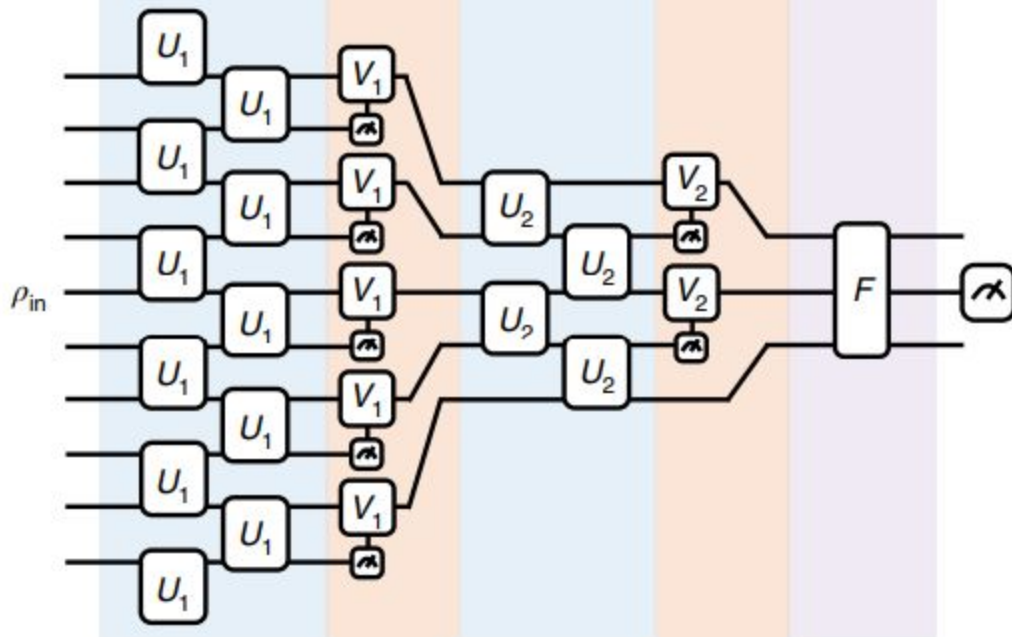
by I Cong · 2019 · Cited by 1206 — We introduce and analyse a **quantum** circuit-based algorithm inspired by **convolutional neural networks**, a highly effective model in machine...

## Scholarly articles for **quantum convolutional neural network**

**Quantum convolutional neural networks** - Cong - Cited by 1206

A tutorial on **quantum convolutional neural networks** ( ... - Oh - Cited by 131

**Quantum convolutional neural network** for classical ... - Hur - Cited by 209

**a****b**

# QCNN

Quantum Convolutional Neural Networks (QCNN) behave in a similar manner to CCNNs.

First, we encode our pixelated image into a quantum circuit using a given feature map, such Qiskit's ZFeatureMap or ZZFeatureMap or others available in the circuit library.

Tabish sagri ~ Hello from my son!

# CNN

After encoding our image, we apply alternating convolutional and pooling layers.

By applying these alternating layers, we reduce the dimensionality of our circuit until we are left with one qubit.

We can then classify our input image by measuring the output of this one remaining qubit.

# QCNN

Quantum Convolutional Neural Networks (QCNN) behave in a similar manner to CCNNs.

First, we encode our pixelated image into a quantum circuit using a given feature map, such Qiskit's ZFeatureMap or ZZFeatureMap or others available in the circuit library.

After encoding our image, we apply alternating convolutional and pooling layers. By applying these alternating layers, we reduce the dimensionality of our circuit until we are left with one qubit. We can then classify our input image by measuring the output of this one remaining qubit.



# QCNN - CONVOLUTION LAYER

The Quantum Convolutional Layer will consist of a series of two qubit unitary operators, which recognize and determine relationships between the qubits in our circuit.

$$U^\dagger U = I_4$$

# QCNN - POOLING LAYER

For the Quantum Pooling Layer, we cannot do the same as is done classically to reduce the dimension, i.e. the number of qubits in our circuit.

Instead, we reduce the number of qubits by performing operations upon each until a specific point and then disregard certain qubits in a specific layer.

It is these layers where we stop performing operations on certain qubits that we call our 'pooling layer'.

# QCNN

In the QCNN, each layer contains parameterized circuits, meaning we alter our output result by adjusting the parameters of each layer.

When training our QCNN, it is these parameters that are adjusted to reduce the loss function of our QCNN.




Figure 2: Example QCNN containing four qubits. The first Convolutional Layer acts on all the qubits. This is followed by the first pooling layer, which reduces the dimensionality of the QCNN from four qubits to two qubits by disregarding the first two. The second Convolutional layer then detects features between the two qubits still in use in the QCNN, followed by another pooling layer, which reduces the dimensionality from two qubits to one, which will be our output qubit.


# QCNN - CONVOLUTION LAYER

```
# We now define a two qubit unitary as defined in [3]
```

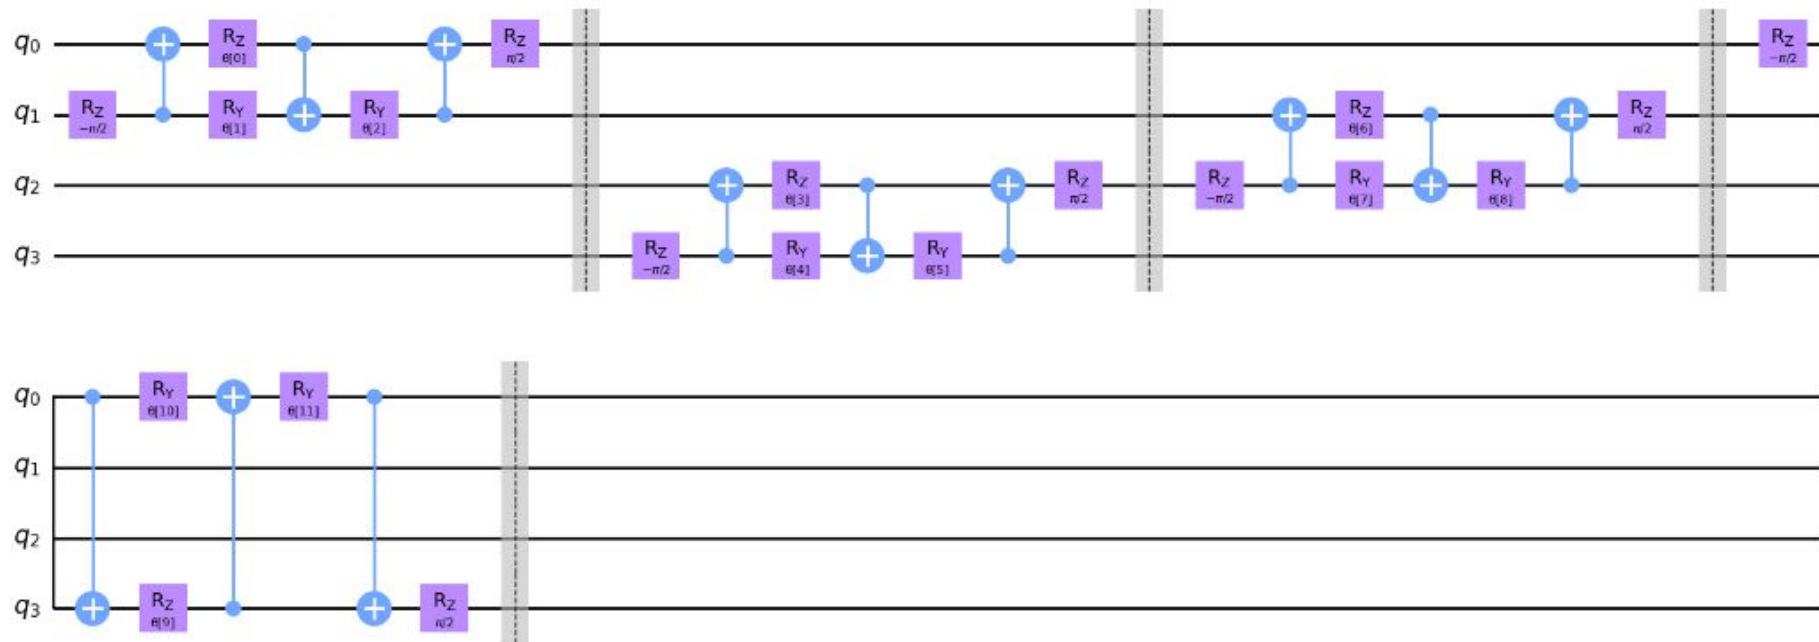
```
def conv_circuit(params):  
    target = QuantumCircuit(2)  
    target.rz(-np.pi / 2, 1)  
    target.cx(1, 0)  
    target.rz(params[0], 0)  
    target.ry(params[1], 1)  
    target.cx(0, 1)  
    target.ry(params[2], 1)  
    target.cx(1, 0)  
    target.rz(np.pi / 2, 0)  
    return target
```

```
# Let's draw this circuit and see what it looks like
```

```
params = ParameterVector("θ", length=3)  
circuit = conv_circuit(params)  
circuit.draw("mpl", style="clifford")
```



# QCNN - CONVOLUTION LAYER



```

: def conv_layer(num_qubits, param_prefix):
    qc = QuantumCircuit(num_qubits, name="Convolutional Layer")
    qubits = list(range(num_qubits))
    param_index = 0
    params = ParameterVector(param_prefix, length=num_qubits * 3)
    for q1, q2 in zip(qubits[0::2], qubits[1::2]):
        qc = qc.compose(conv_circuit(params[param_index : (param_index + 3)]), [q1, q2])
        qc.barrier()
        param_index += 3
    for q1, q2 in zip(qubits[1::2], qubits[2::2] + [0]):
        qc = qc.compose(conv_circuit(params[param_index : (param_index + 3)]), [q1, q2])
        qc.barrier()
        param_index += 3

    qc_inst = qc.to_instruction()

    qc = QuantumCircuit(num_qubits)
    qc.append(qc_inst, qubits)
    return qc

```

```

circuit = conv_layer(4, "θ")
circuit.decompose().draw("mpl", style="clifford")

```

# POOLING LAYER

The purpose of a pooling layer is to reduce the dimensions of our Quantum Circuit, i.e. reduce the number of qubits in our circuit, while retaining as much information as possible from previously learned data.


Reducing the amount of qubits also reduces the computational cost of the overall circuit, as the number of parameters that the QCNN needs to learn decreases.




# POOLING LAYER

```
def pool_circuit(params):  
    target = QuantumCircuit(2)  
    target.rz(-np.pi / 2, 1)  
    target.cx(1, 0)  
    target.rz(params[0], 0)  
    target.ry(params[1], 1)  
    target.cx(0, 1)  
    target.ry(params[2], 1)  
  
    return target
```

```
params = ParameterVector("θ", length=3)  
circuit = pool_circuit(params)  
circuit.draw("mpl", style="clifford")
```



# POOLING LAYER



```

: def pool_layer(sources, sinks, param_prefix):
    num_qubits = len(sources) + len(sinks)
    qc = QuantumCircuit(num_qubits, name="Pooling Layer")
    param_index = 0
    params = ParameterVector(param_prefix, length=num_qubits // 2 * 3)
    for source, sink in zip(sources, sinks):
        qc = qc.compose(pool_circuit(params[param_index : (param_index + 3)]), [source, sink])
        qc.barrier()
        param_index += 3





    qc_inst = qc.to_instruction()

    qc = QuantumCircuit(num_qubits)
    qc.append(qc_inst, range(num_qubits))
    return qc

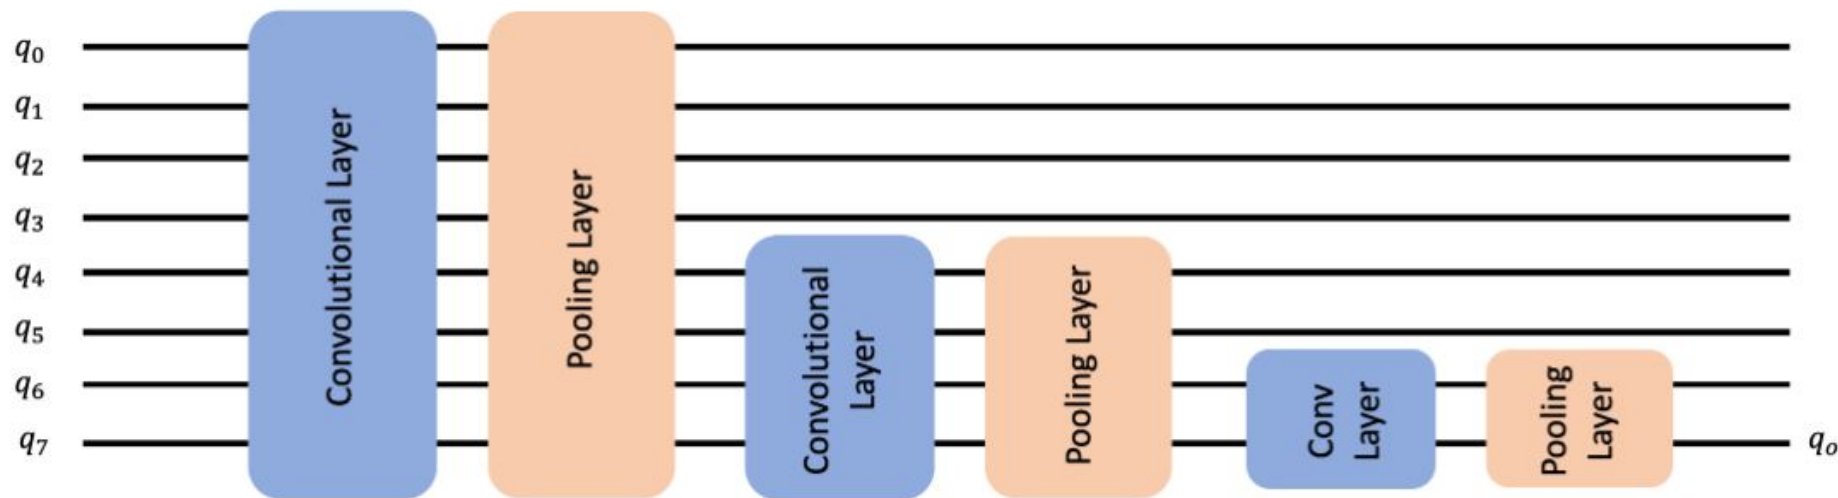
sources = [0, 1]
sinks = [2, 3]
circuit = pool_layer(sources, sinks, "θ")
circuit.decompose().draw("mpl", style="clifford")

```

# EXAMPLE



# EXAMPLE



```
: feature_map = ZFeatureMap(8)

ansatz = QuantumCircuit(8, name="Ansatz")

# First Convolutional Layer
ansatz.compose(conv_layer(8, "c1"), list(range(8)), inplace=True)

# First Pooling Layer
ansatz.compose(pool_layer([0, 1, 2, 3], [4, 5, 6, 7], "p1"), list(range(8)), inplace=True)



# Second Convolutional Layer
ansatz.compose(conv_layer(4, "c2"), list(range(4, 8)), inplace=True)

# Second Pooling Layer
ansatz.compose(pool_layer([0, 1], [2, 3], "p2"), list(range(4, 8)), inplace=True)

# Third Convolutional Layer
ansatz.compose(conv_layer(2, "c3"), list(range(6, 8)), inplace=True)

# Third Pooling Layer
ansatz.compose(pool_layer([0], [1], "p3"), list(range(6, 8)), inplace=True)

# Combining the feature map and ansatz
circuit = QuantumCircuit(8)
circuit.compose(feature_map, range(8), inplace=True)
circuit.compose(ansatz, range(8), inplace=True)
```



# REFERENCES FOR QCNN

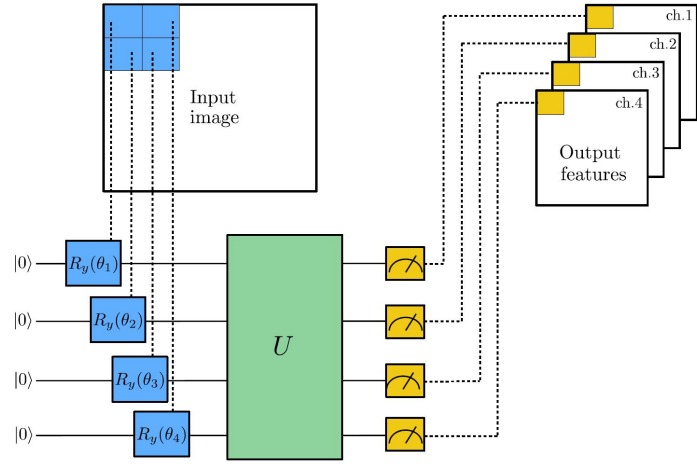
[https://qiskit-community.github.io/qiskit-machine-learning/tutorials/11 quantum convolutional neural networks.html](https://qiskit-community.github.io/qiskit-machine-learning/tutorials/11_quantum_convolutional_neural_networks.html)

<https://www.ibm.com/topics/convolutional-neural-networks>

<https://sci-hub.se/https://www.nature.com/articles/s41567-019-0648-8>



# QUANVOLUTIONAL NEURAL NETWORKS



# **Quanvolutional Neural Networks: Powering Image Recognition with Quantum Circuits**

Maxwell Henderson<sup>1 \*</sup>, Samriddhi Shakya<sup>1</sup>, Shashindra Pradhan<sup>1</sup>, and Tristan Cook<sup>1</sup>

<sup>1</sup>*QxBranch, Inc., 777 6th St NW, 11th Floor, Washington DC, 20001*

Convolutional neural networks (CNNs) have rapidly risen in popularity for many machine learning applications, particularly in the field of image recognition. Much of the benefit generated from these networks comes from their ability to extract features from the data in a hierarchical manner. These features are extracted using various transformational layers, notably the convolutional layer which gives the model its name. In this work, we introduce a new type of transformational layer called a quantum convolution, or quanvolutional layer. Quanvolutional layers operate on input data by locally transforming the data using a number of random quantum circuits, in a way that is similar to the transformations performed by random convolutional filter layers. Provided these quantum transformations produce meaningful features for classification purposes, then the overall algorithm could be quite useful for near term quantum computing, because it requires small quantum circuits with little to no error correction. In this work, we empirically evaluated the potential benefit of these quantum transformations by comparing three types of models built on the MNIST dataset: CNNs, quantum convolutional neural networks (QNNs), and CNNs with additional non-linearities introduced. Our results showed that the QNN models had both higher test set accuracy as well as faster training compared to the purely classical CNNs.

# QUANVOLUTIONAL

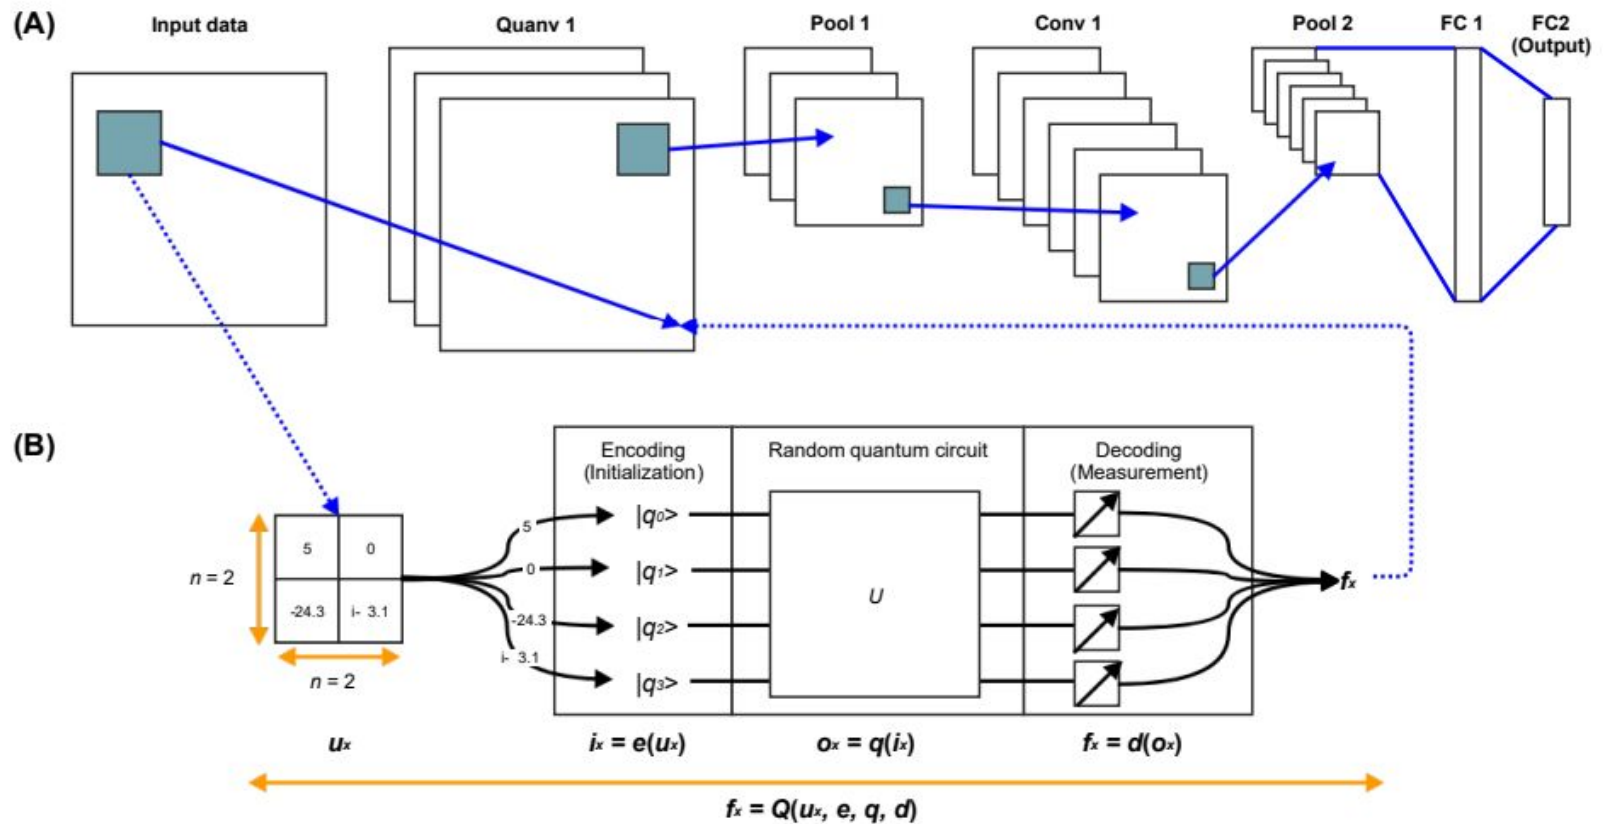
Quanvolutional Layers enhance or replace traditional convolutional layers in CNNs. They operate as follows:

**Quantum Circuit Embedding:** Input data (e.g., image patches) is encoded into a quantum state using a quantum circuit.

**Quantum Transformation:** The quantum state undergoes transformations via quantum gates, exploiting properties like superposition and entanglement to extract complex features.

**Measurement:** The quantum state is measured to produce classical data, which serves as features for the neural network.

**Feature Extraction:** The classical data from quantum measurements is used as features in the neural network, passed to subsequent layers for further processing and classification.



**Fig. 1.:** A. Simple example of a quannvolutional layer in a full network stack. The quannvolutional layer contains several quannvolutional filters (three in this example) that transform the input data into different output feature maps. B. An in-depth look at the processing of classical data into and out of the random quantum circuit in the quannvolutional filter.



# EQUATIONS FOR A QUANVOLUTIONAL LAYER

$$i_x = e(u_x).$$

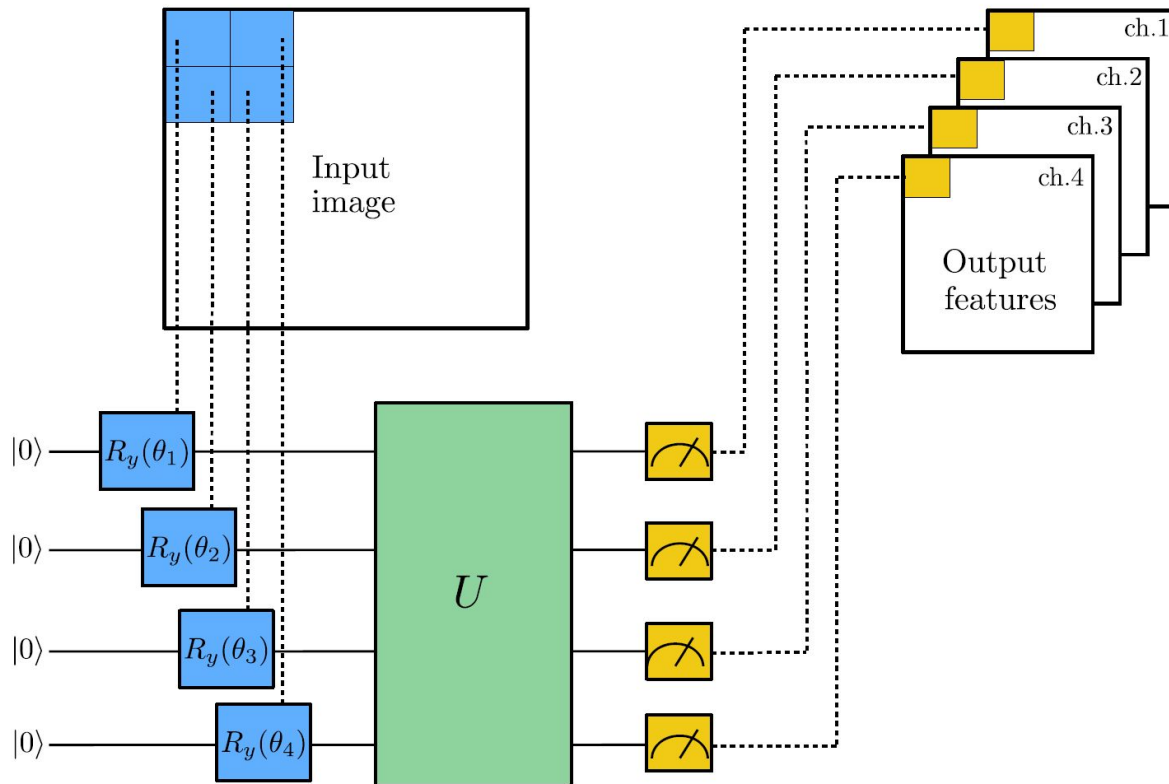
$$o_x = q(i_x) = q(e(u_x)).$$

$$f_x = d(o_x) = d(q(e(u_x)))$$

$$f_x = Q(u_x, e, q, d).$$



# QUANVOLUTIONAL NEURAL NETWORK




# QUANVOLUTIONAL LAYER


1. A small region of the input image, in our example a  $2 \times 2$  square, is embedded into a quantum circuit. In this demo, this is achieved with parametrized rotations applied to the qubits initialized in the ground state.
2. A quantum computation, associated to a unitary  $U$  is performed on the system. The unitary could be generated by a variational quantum circuit or, more simply, by a random circuit as proposed .
3. The quantum system is finally measured, obtaining a list of classical expectation values. The measurement results could also be classically post-processed but, for simplicity, in this demo we directly use the raw expectation values.
4. Analogously to a classical convolution layer, each expectation value is mapped to a different channel of a single output pixel.
5. Iterating the same procedure over different regions, one can scan the full input image, producing an output object which will be structured as a multi-channel image.
6. The quantum convolution can be followed by further quantum layers or by classical layers.



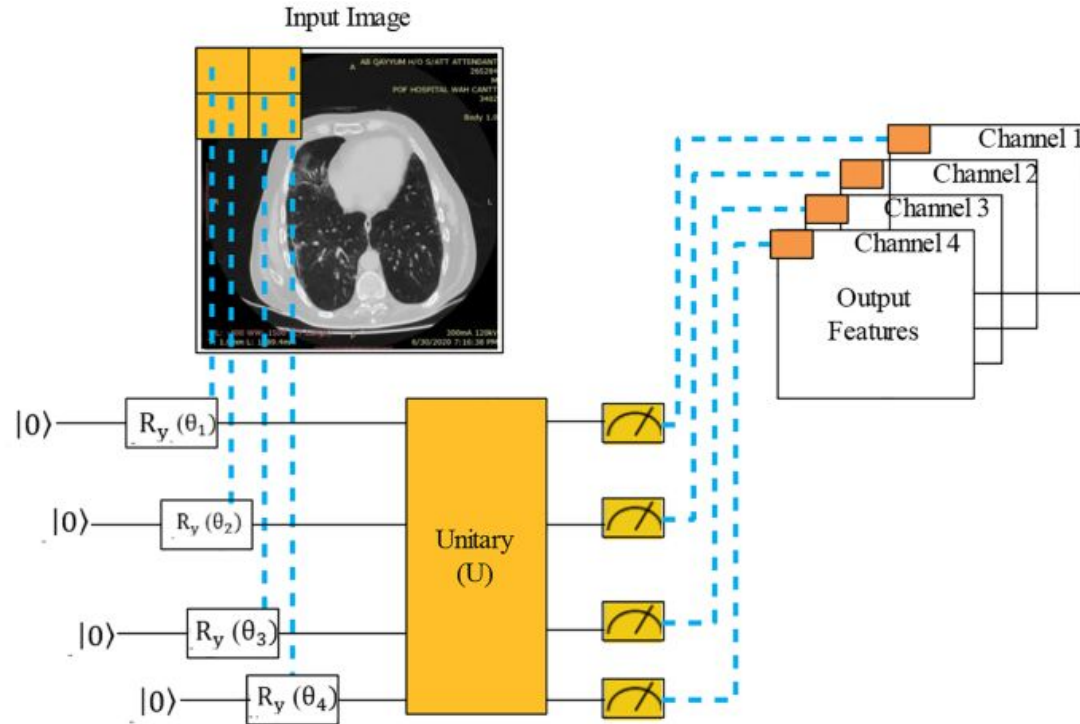
# RESULTS GENERATED BY PENNYLANE



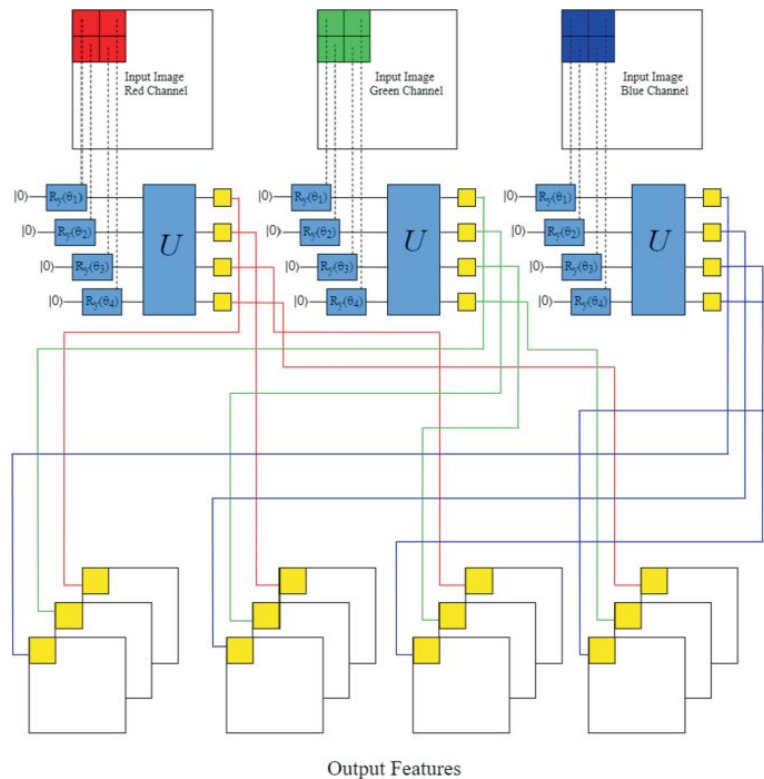
# RESULTS GENERATED BY PENNYLANE



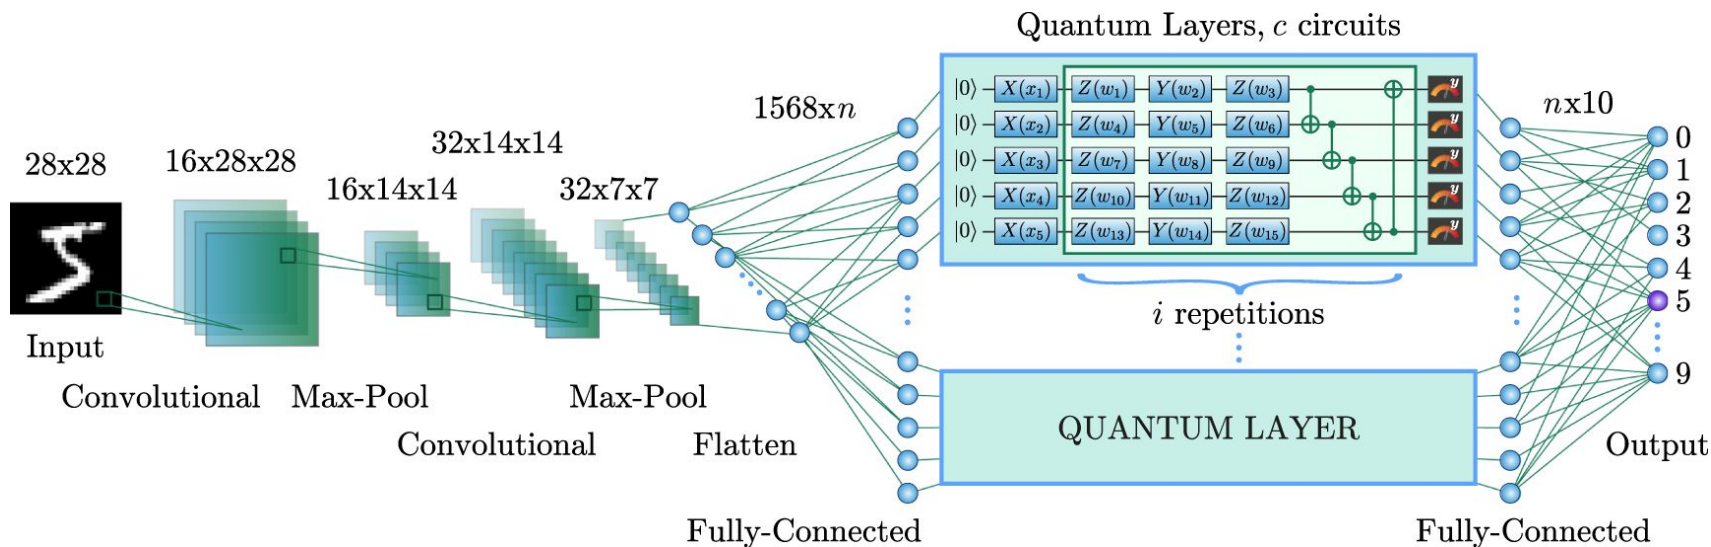
# QUANTUM MACHINE LEARNING ARCHITECTURE FOR COVID-19 CLASSIFICATION BASED ON SYNTHETIC DATA GENERATION



# GQNN: GREEDY QUANTVOLUTIONAL NEURAL NETWORK MODEL



# QUANTUM MACHINE LEARNING FOR IMAGE CLASSIFICATION



# REFERENCES FOR QCNN

<https://arxiv.org/pdf/1904.04767>

[https://pennylane.ai/qml/demos/tutorial\\_quanvolution/](https://pennylane.ai/qml/demos/tutorial_quanvolution/)

<https://github.com/viventriglia/Quantum Neural Network QNN>

<https://link.springer.com/chapter/10.1007/978-3-031-12413-6>

31